

# The Big Picture

Thierry Sans

# Goals of this lecture

- Define what an Operating System is
- Explain how an OS works in a nutshell
- Bridge the gap between hardware (CSCB58) and systems programming (CSCB09)
- Give an overview of the course content and projects


# The big picture in 5 pieces

The need for <b>bootstrapping</b> and <b>system calls</b>	project 0
The need for <b>concurrency</b>	project 1
The need for <b>user spaces</b>	project 2
The need for <b>virtual memory</b>	project 3
The need for <b>a filesystem</b>	project 4

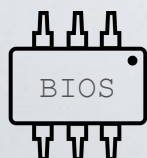
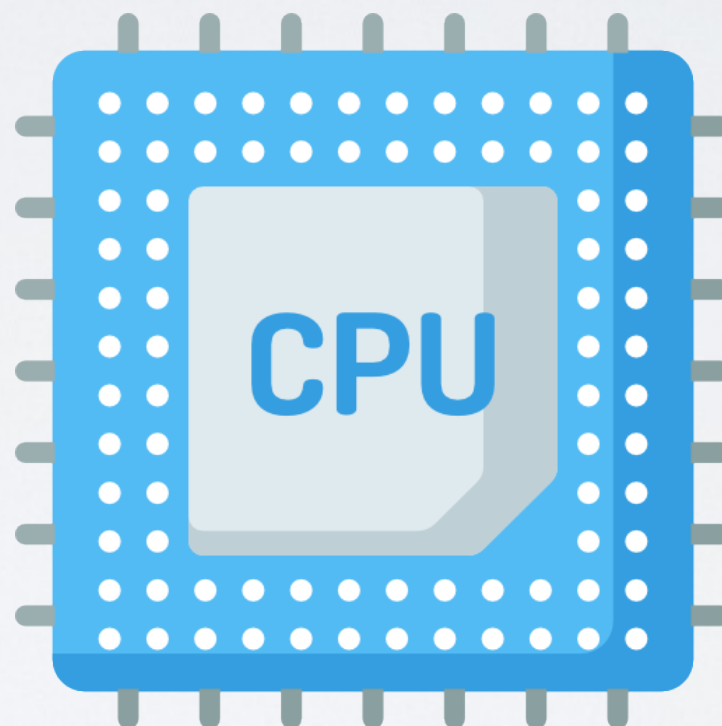
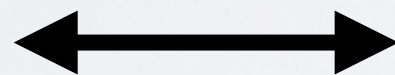
0x FF FF FF FF



I/O



RAM



Boot

0x 00 00 00 00

# Simple Computer Architecture

## Memory + CPU

for a more accurate and detailed map of the x86 memory  
look at [https://wiki.osdev.org/Memory\\_Map\\_\(x86\)](https://wiki.osdev.org/Memory_Map_(x86))

# Each processor has its Instruction Set Architecture (ISA)

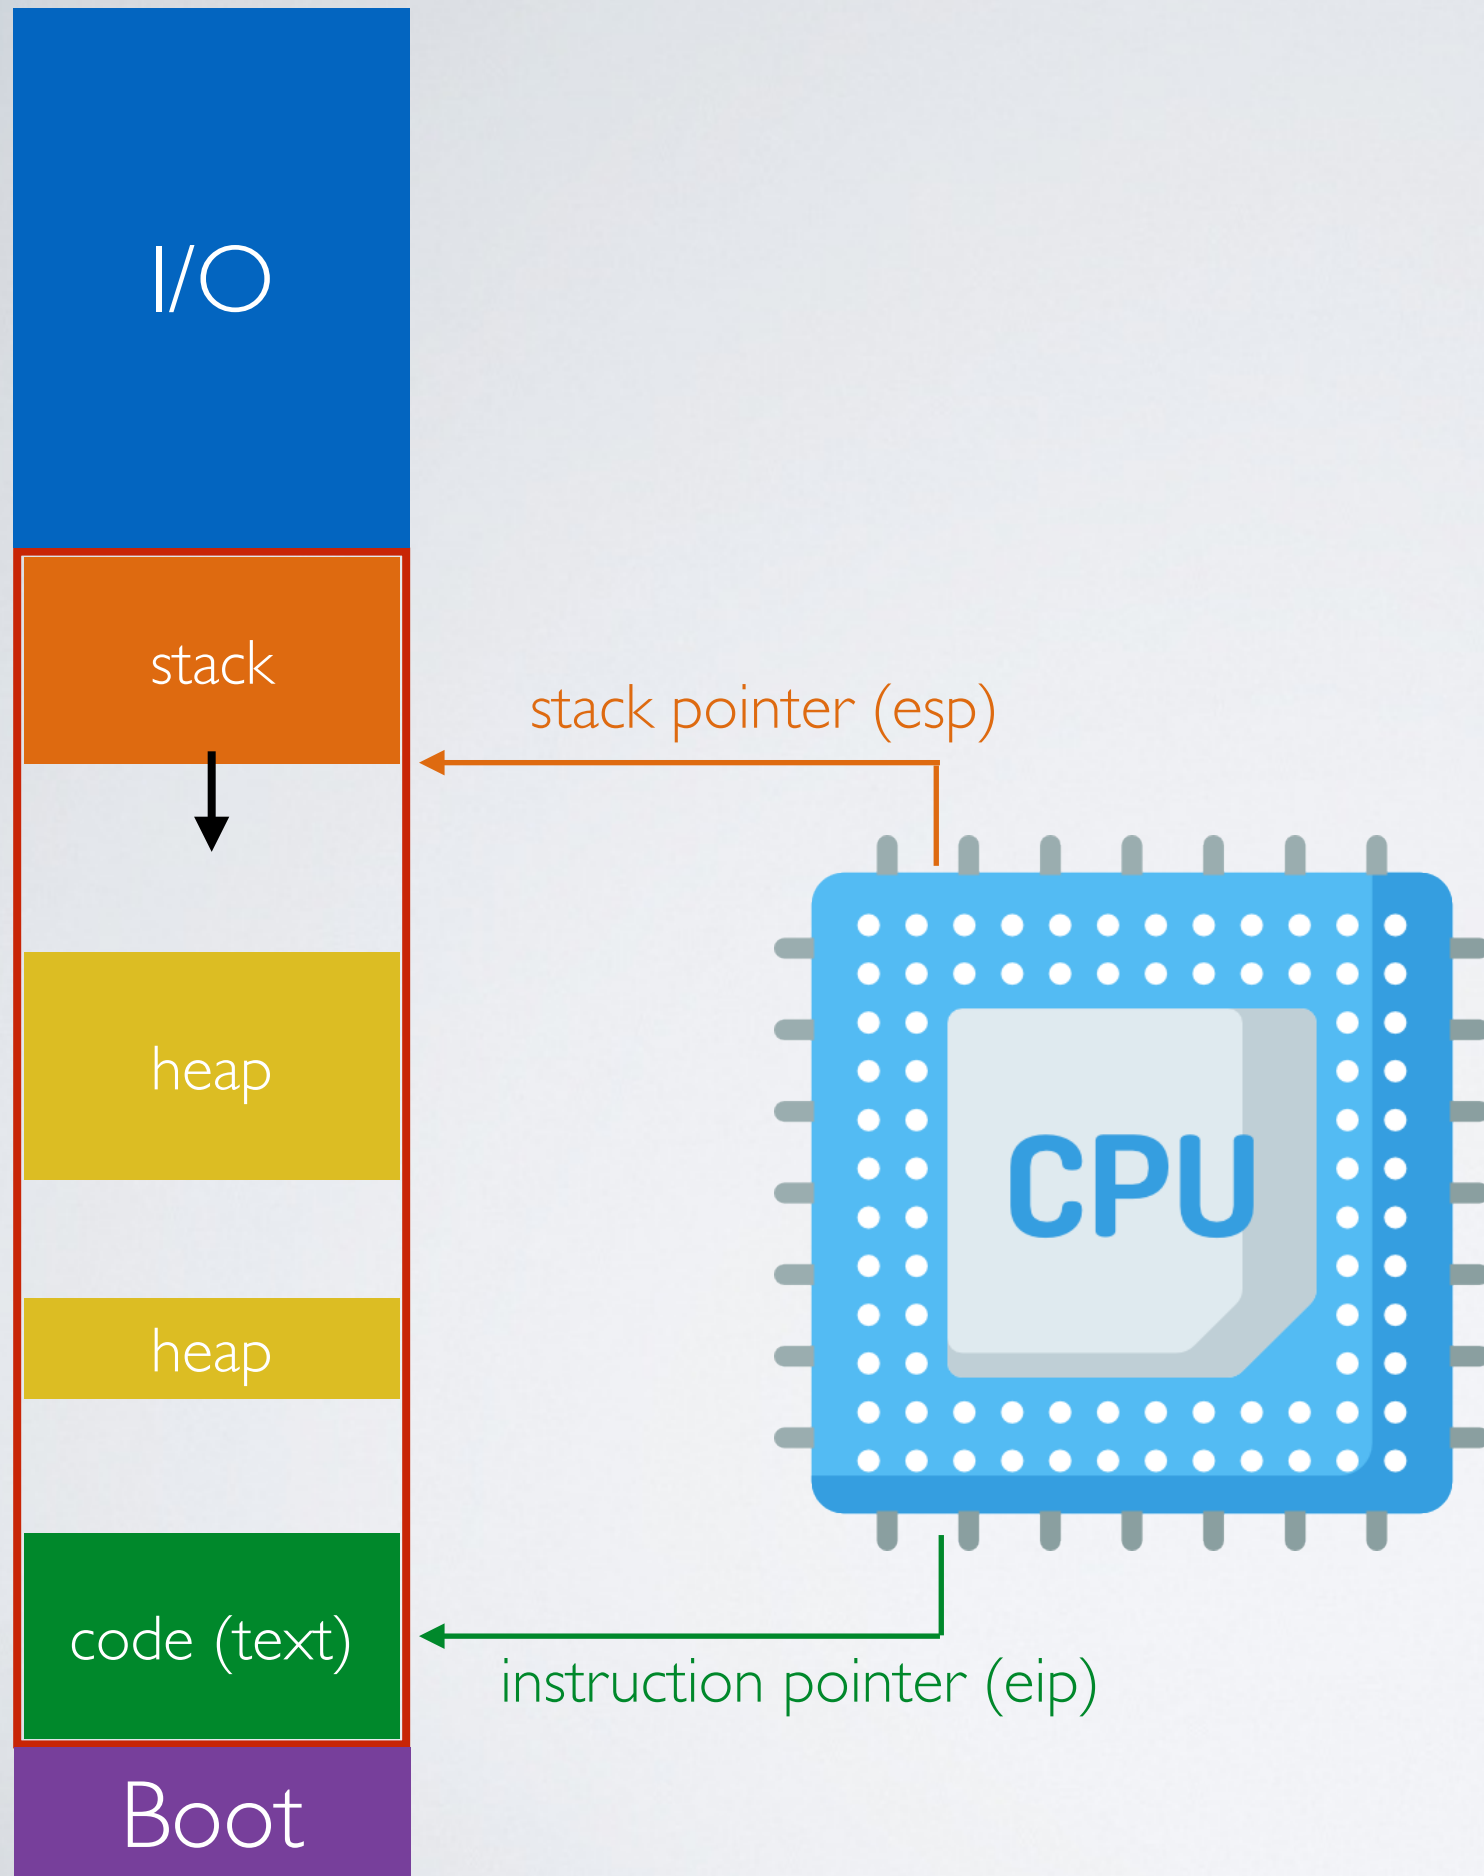
Processor executes instructions stored in memory

- ➡ Each instruction is a bit string that the processor understands as an operation
  - arithmetic
  - read/write bit strings
  - bit logic
  - jumps
- ✓ ~2000 instructions on modern x86-64 processors

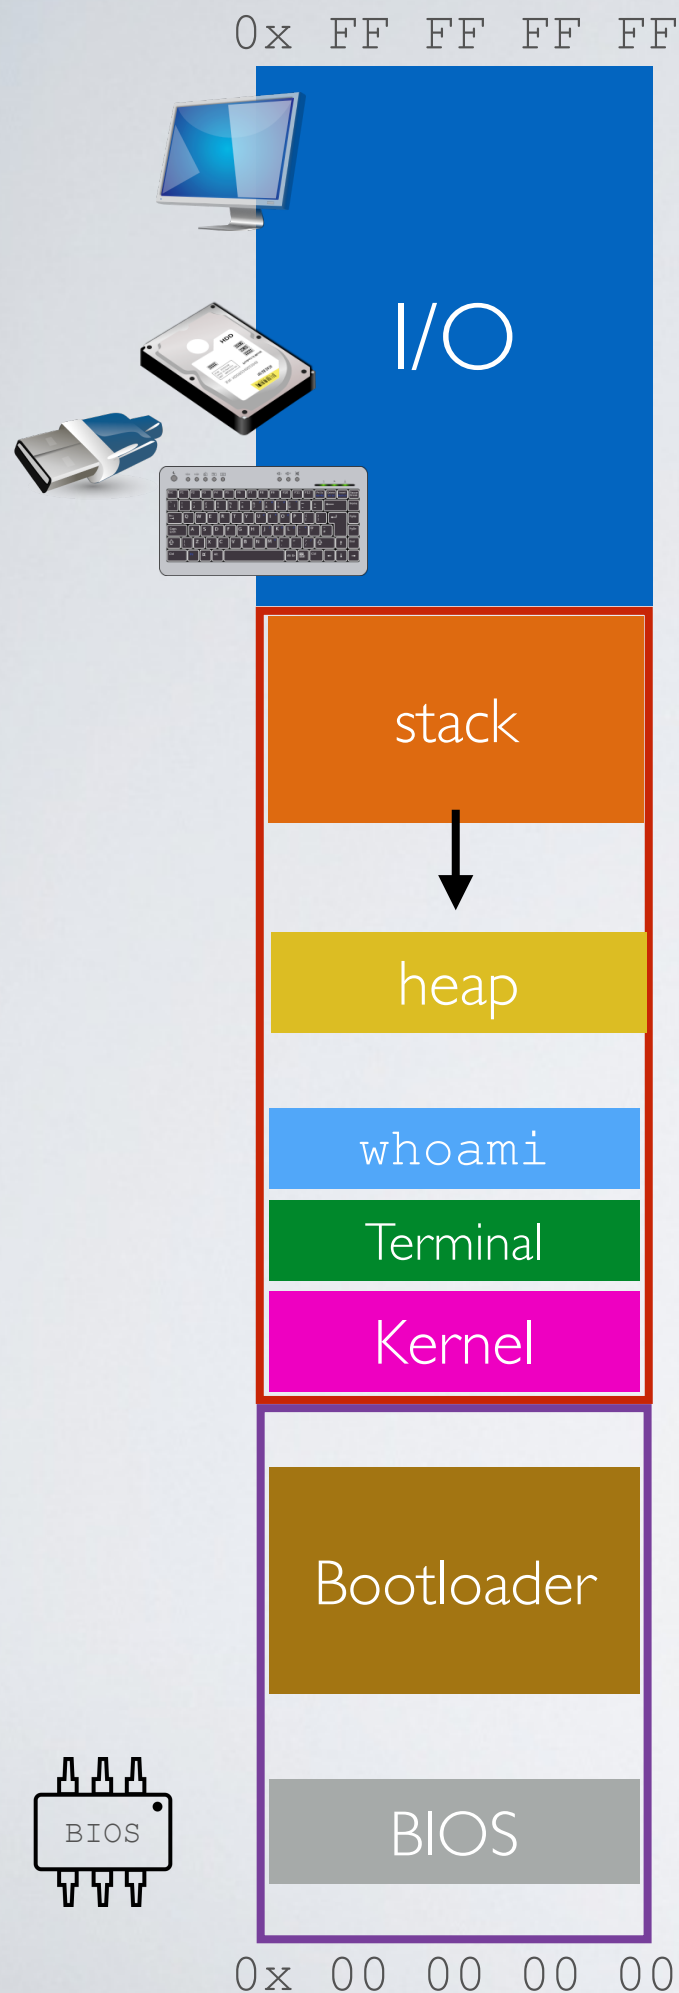


0x FF FF FF FF

Running one program



The need for **bootstrapping**  
and **system calls**



# Bootstrapping

Step 5: using the terminal, users can execute programs (e.g Bash terminal) ... and repeat

Step 4: the kernel starts the user-interface program (e.g Bash terminal)

Step 3: the bootloader loads the OS kernel in RAM

Step 2: the BIOS loads the **bootloader** from a device (hard-drive, USB, network ...) based on the configuration

Step 1: Power -on! The CPU starts executing code contained in the **BIOS** (basic input/output system)



# The need for abstraction for user programs

How to write a user program like the *Bash* shell that reads keyboard inputs from the user?

➡ Read input data from the I/O device directly? But which one?

- The one connected to the PS2 port?
- The one connected to the USB?
- The one connected to the bluetooth?
- The remote one connected to the network?

⦿ User programs do not operate I/O devices directly

✓ The OS abstracts those functionalities and provide them as **system calls**

# System Calls

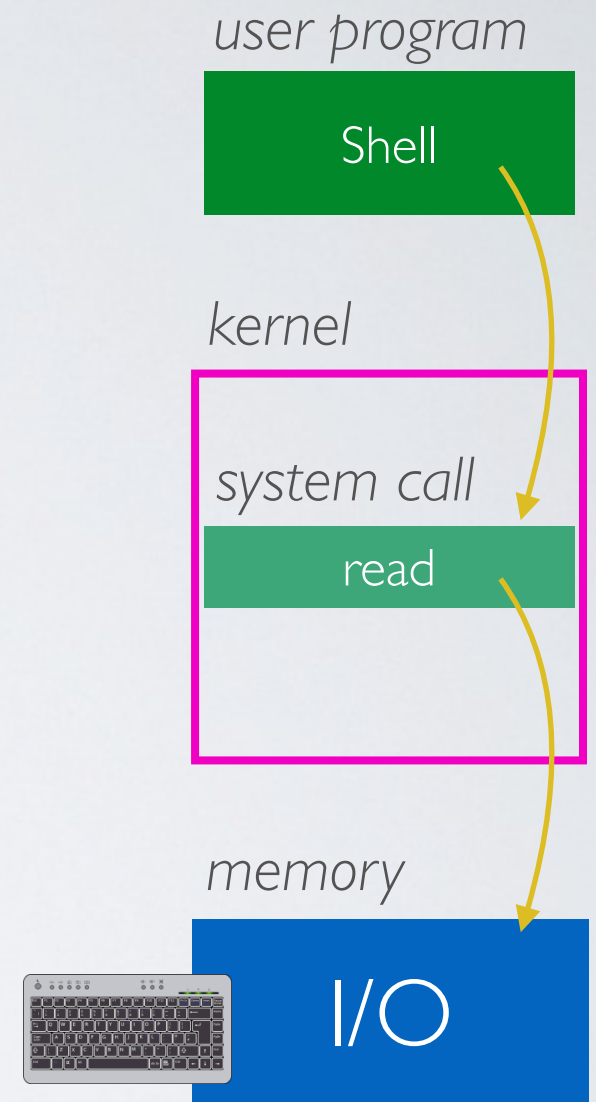
➔ Provide user programs with an API to use the services of operating system

There are 5 categories of system calls

- Process control
- File management
- Device management
- Information/maintenance (system configuration)
- Communication (IPC)
- Protection

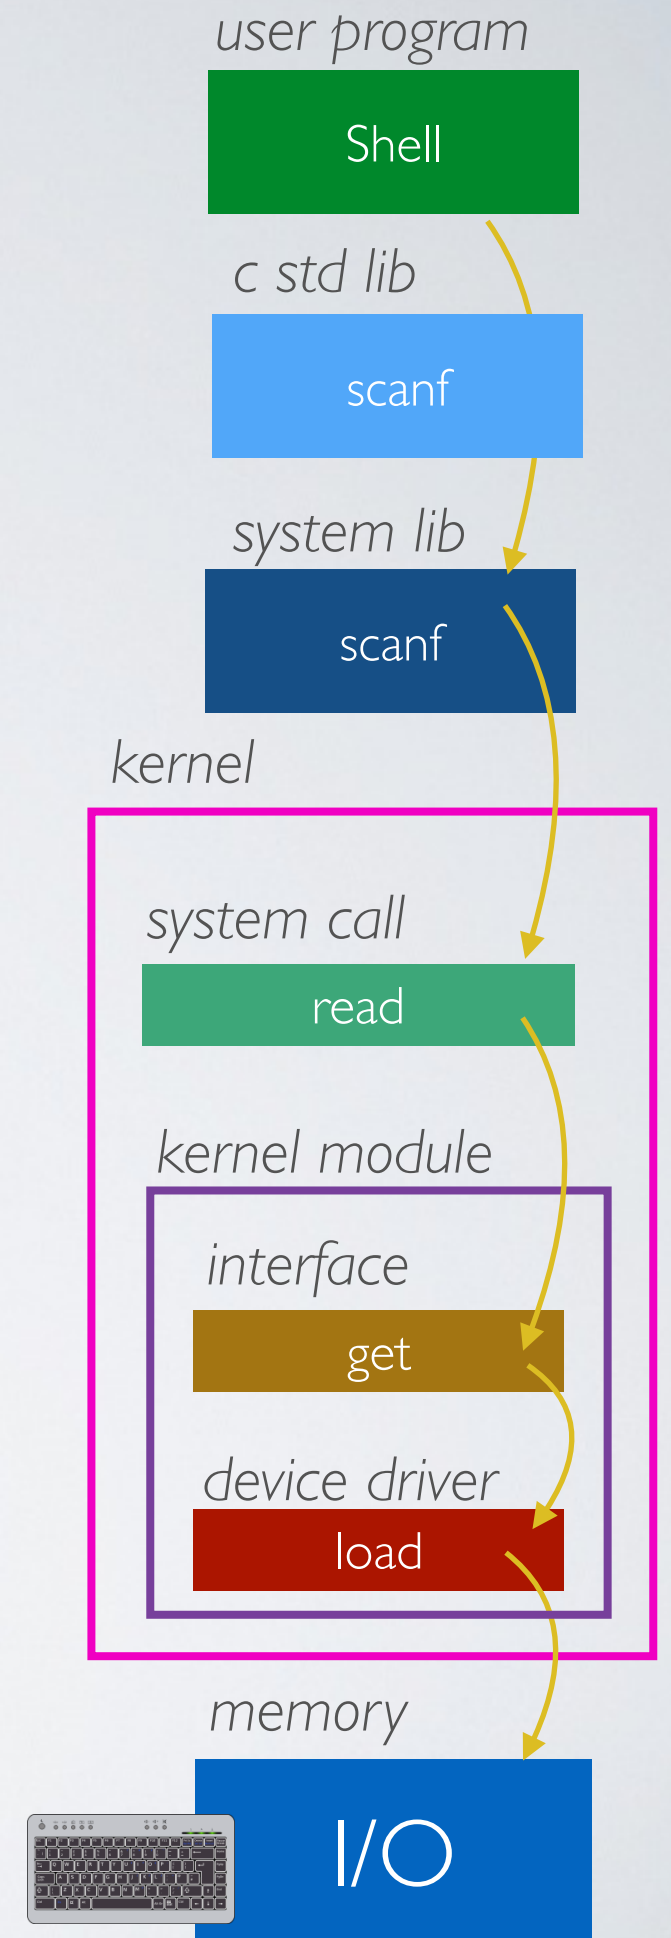
✓ There are 393 system calls on Linux 3.7

[http://www.cheat-sheets.org/saved-copy/Linux\\_Syscall\\_quickref.pdf](http://www.cheat-sheets.org/saved-copy/Linux_Syscall_quickref.pdf)



In reality, many (many) level of abstraction and modularity

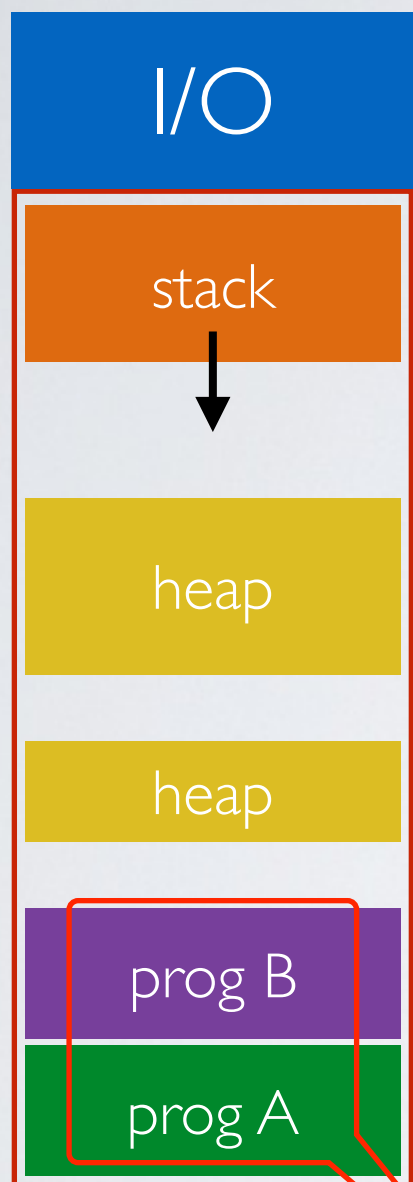
➔ This is what makes developing OS very challenging (CSCB07)




The need for **concurrency**



# Running multiple programs one after the other

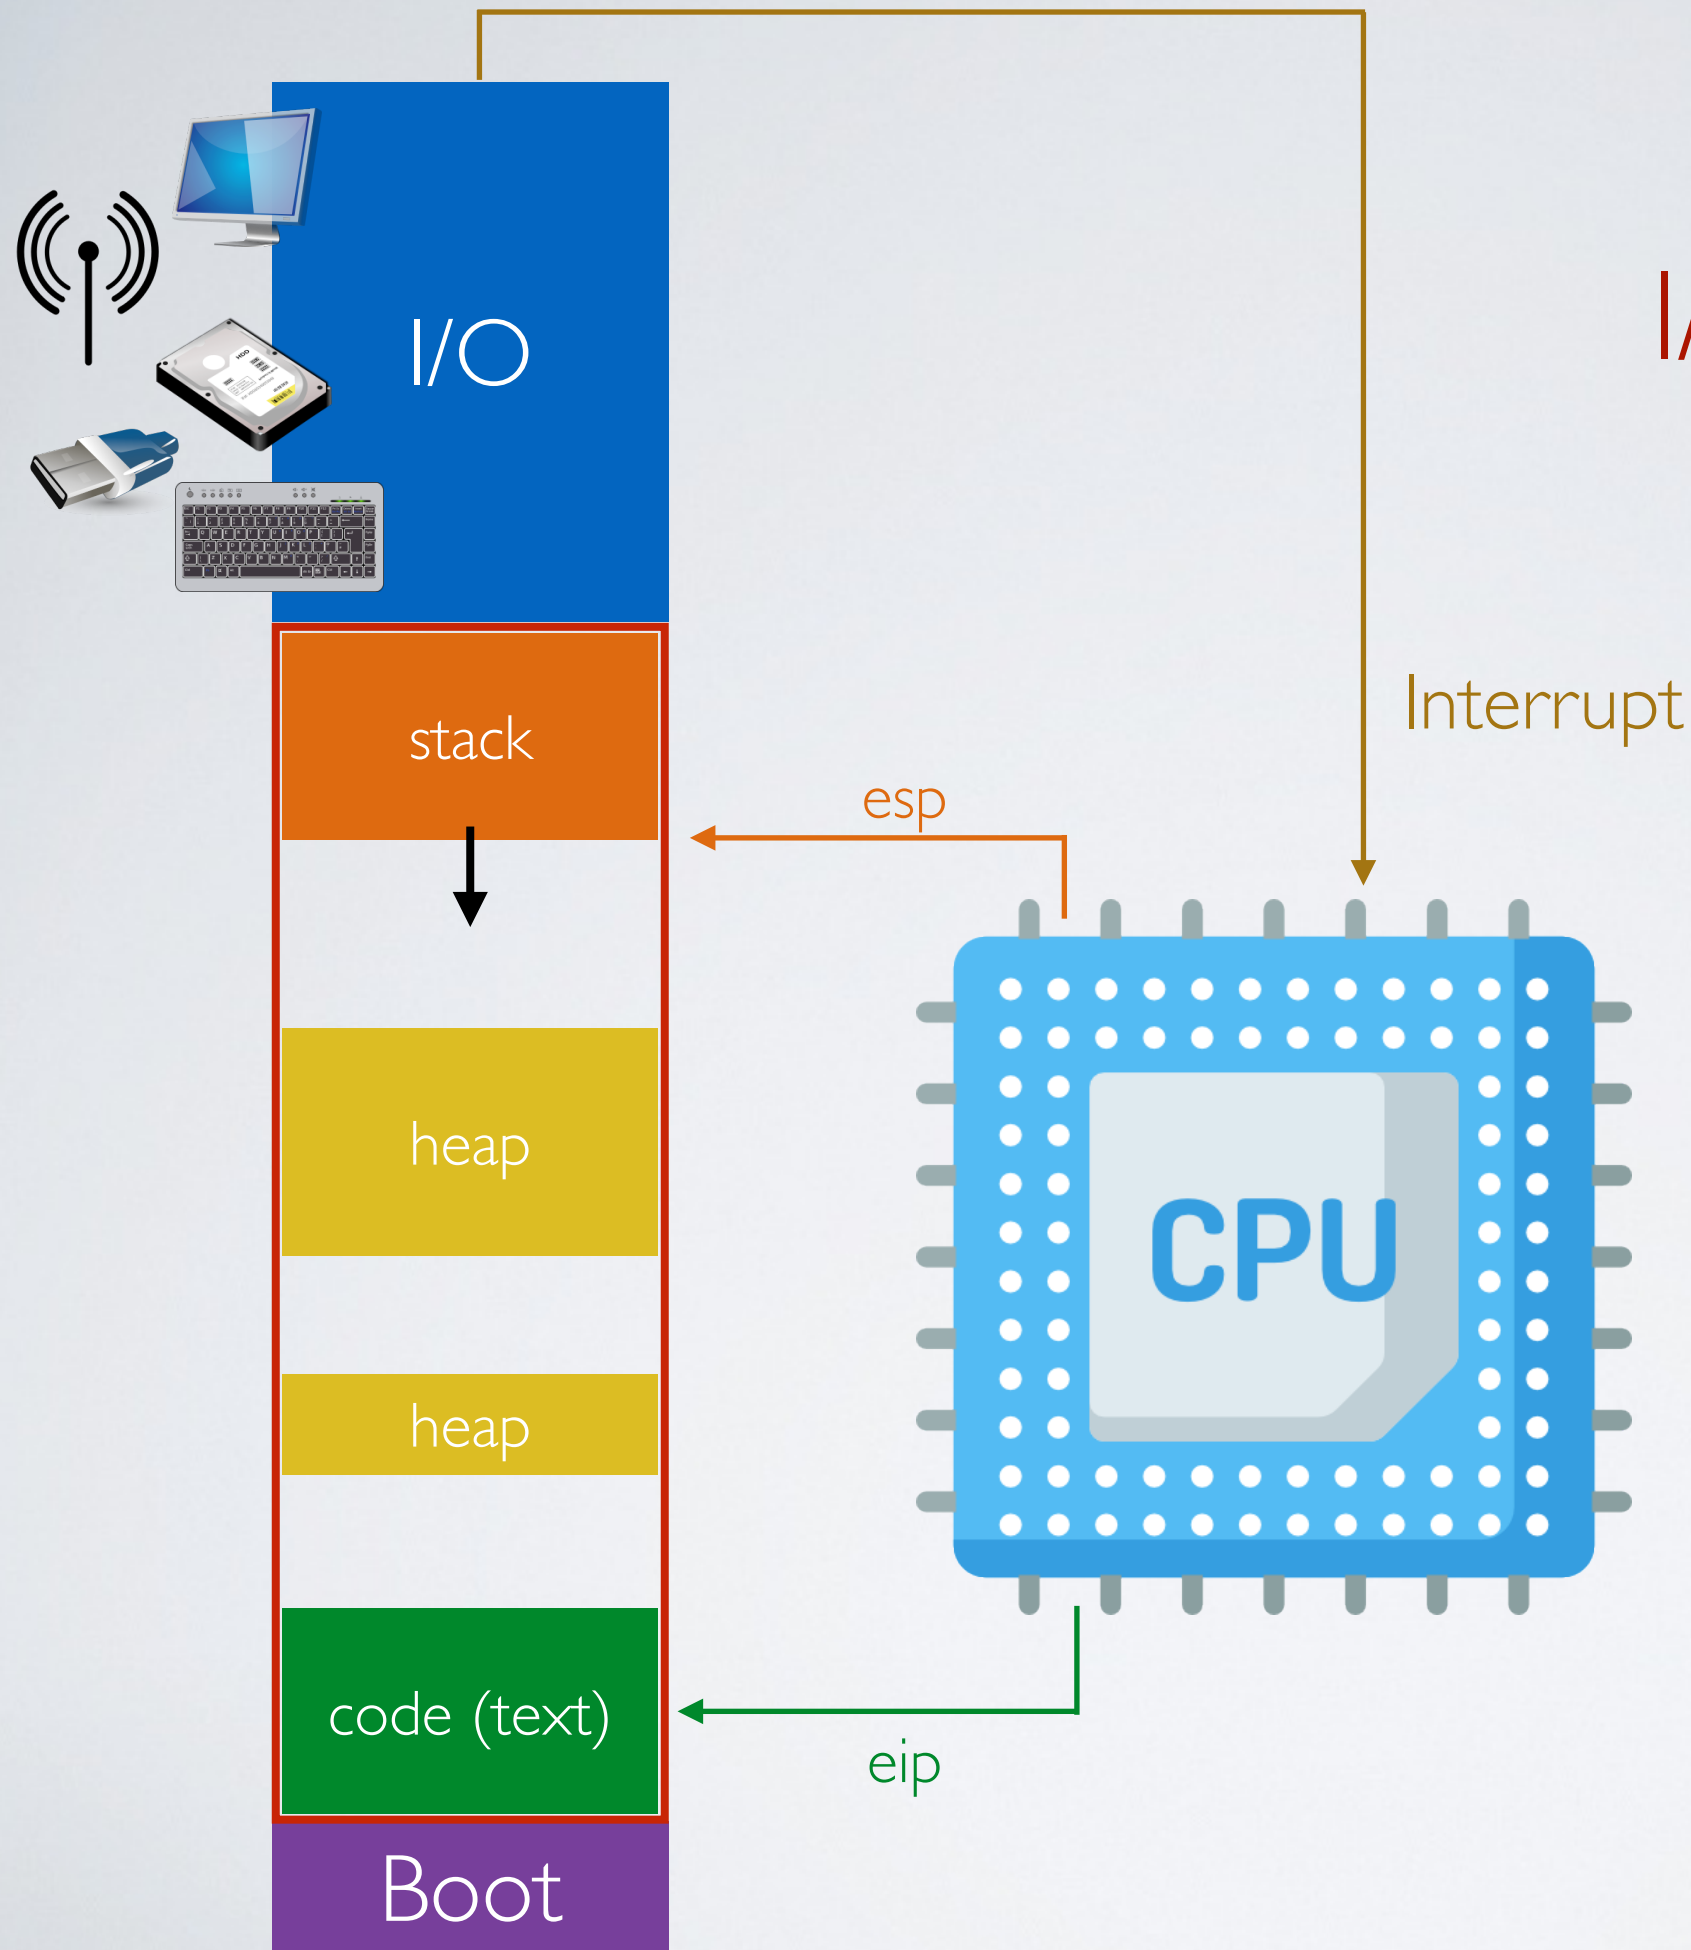


**Problem:** the CPU is waiting for I/O (polling)



**Problem:** the programs must co-exists in memory  
(coming next with virtual memory)

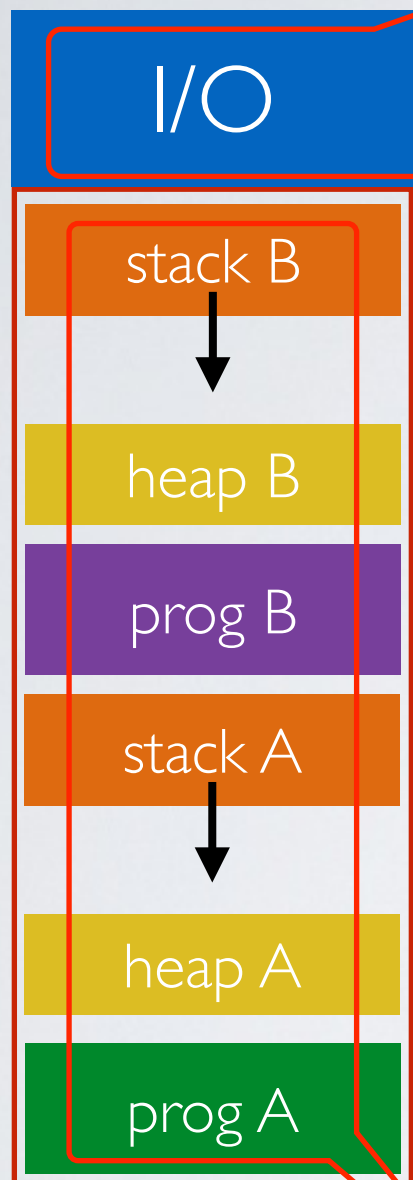





I/O with interrupts

# Running multiple programs concurrently

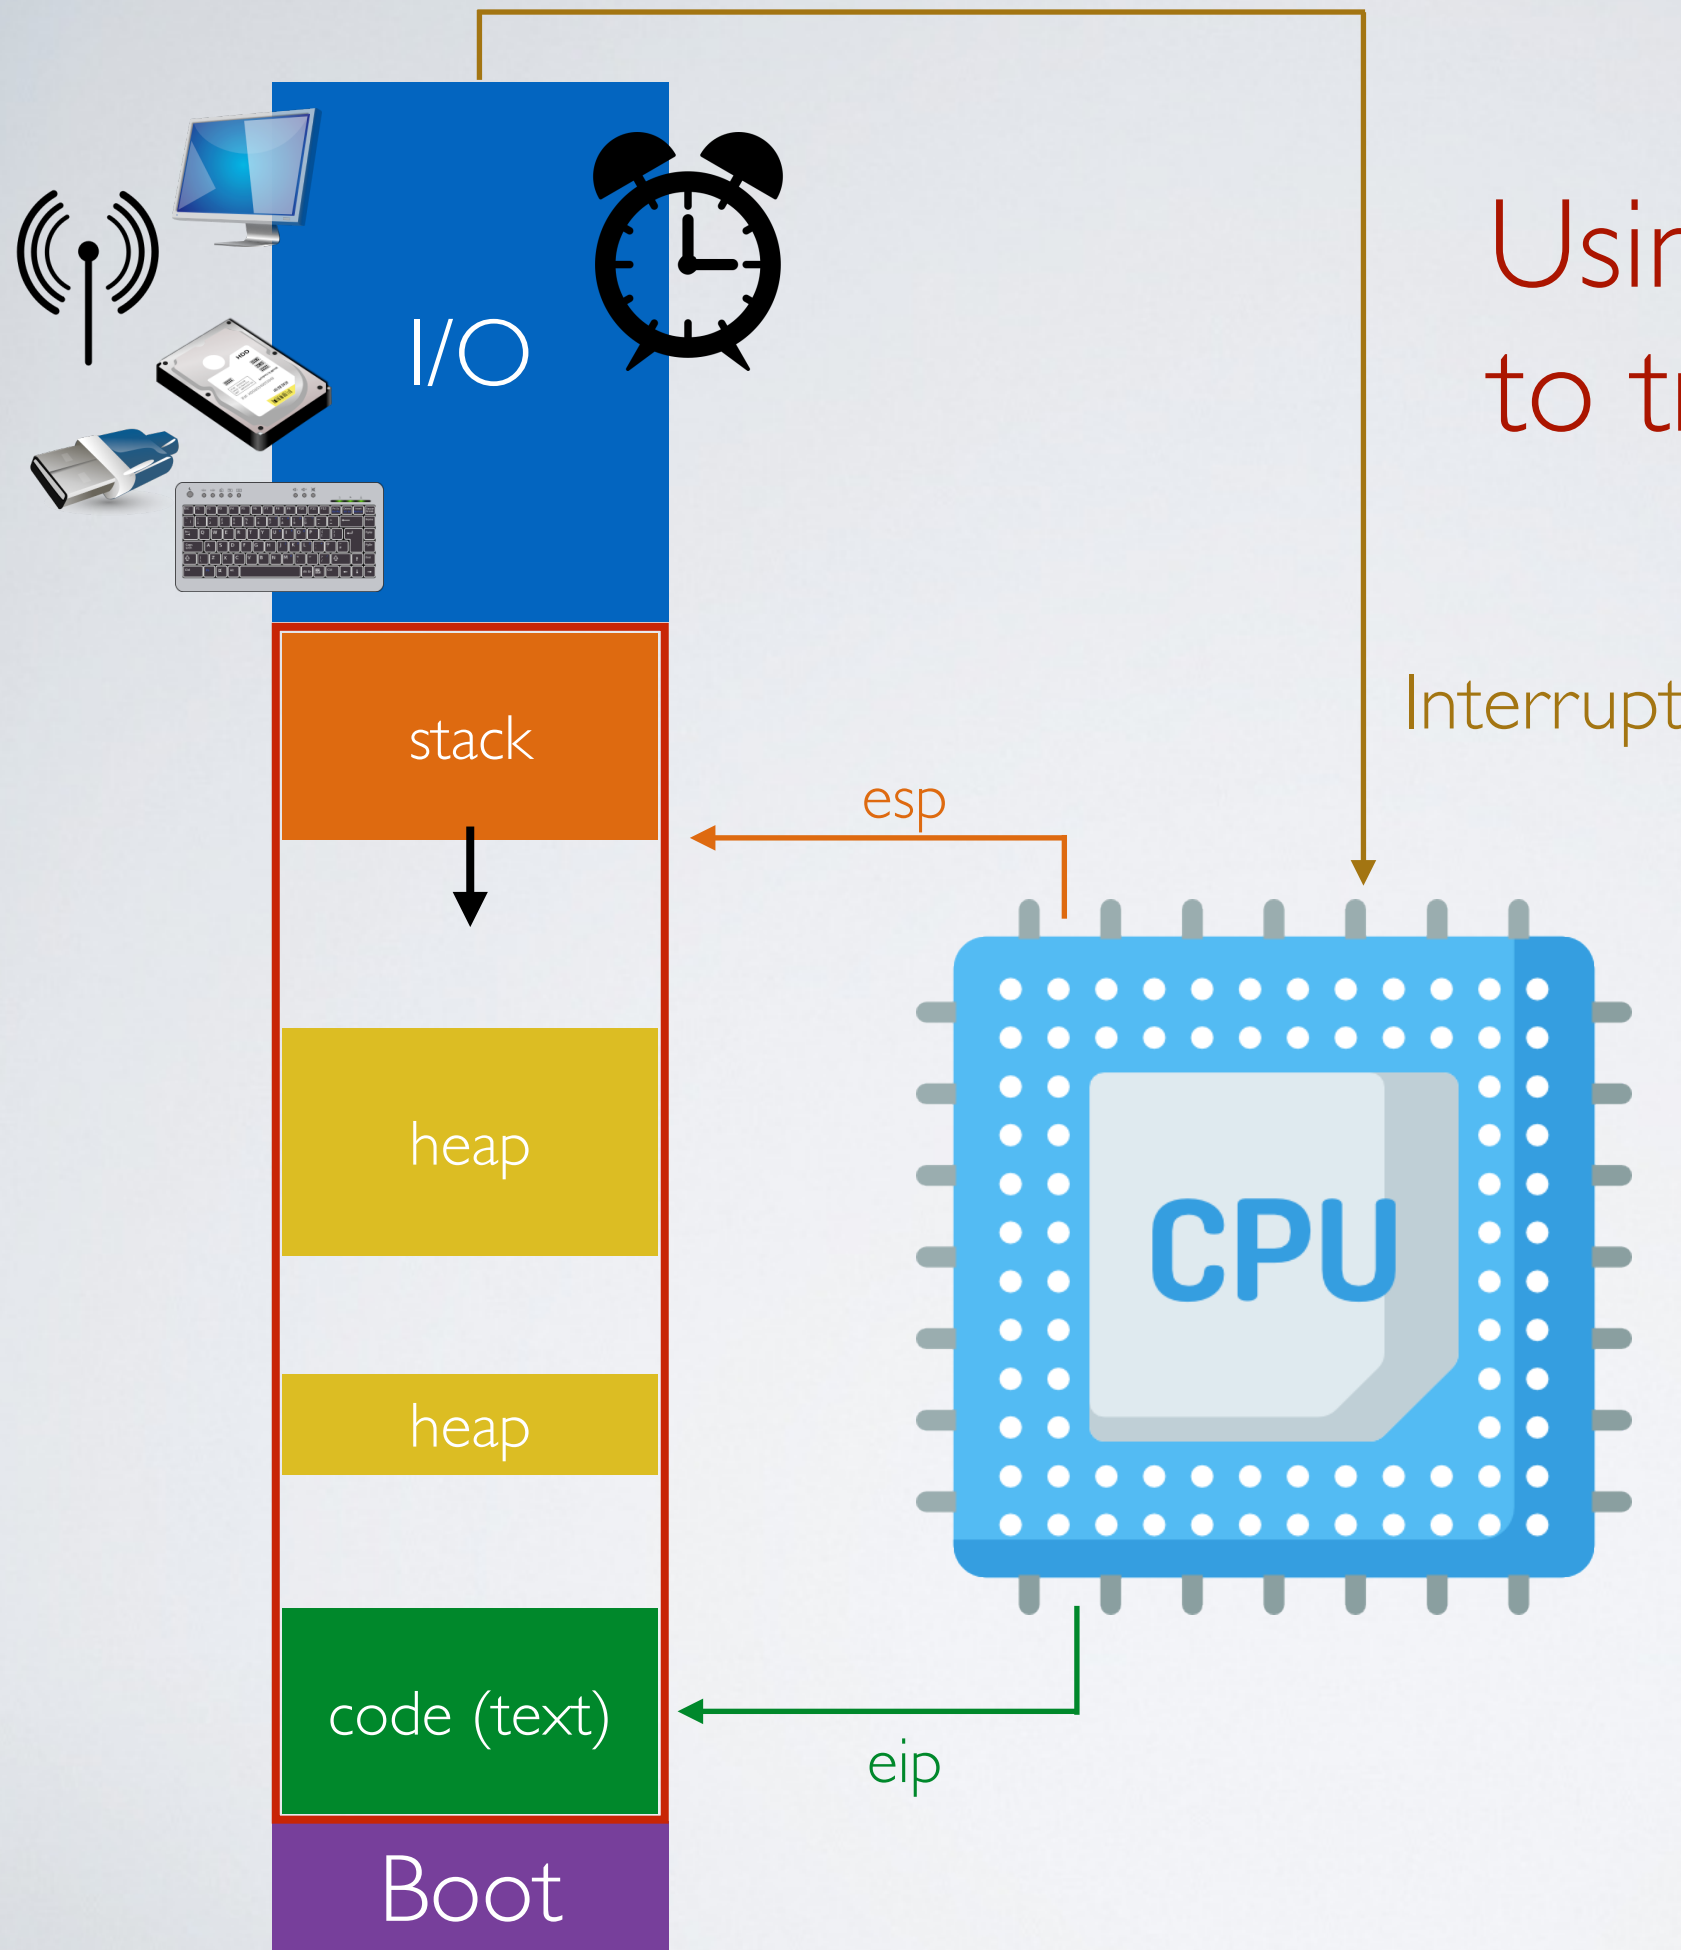
**Problem:** concurrent access to I/O devices must be synchronized



**Problem:** what if the program does not do any I/O and use the CPU for a long time (a.k.a starvation problem)




**Problem:** the programs and their stacks must co-exists in memory (coming next with virtual memory)



Using the clock  
to trigger an interrupt

# Process States





# With concurrency

- ✓ **From the system perspective**

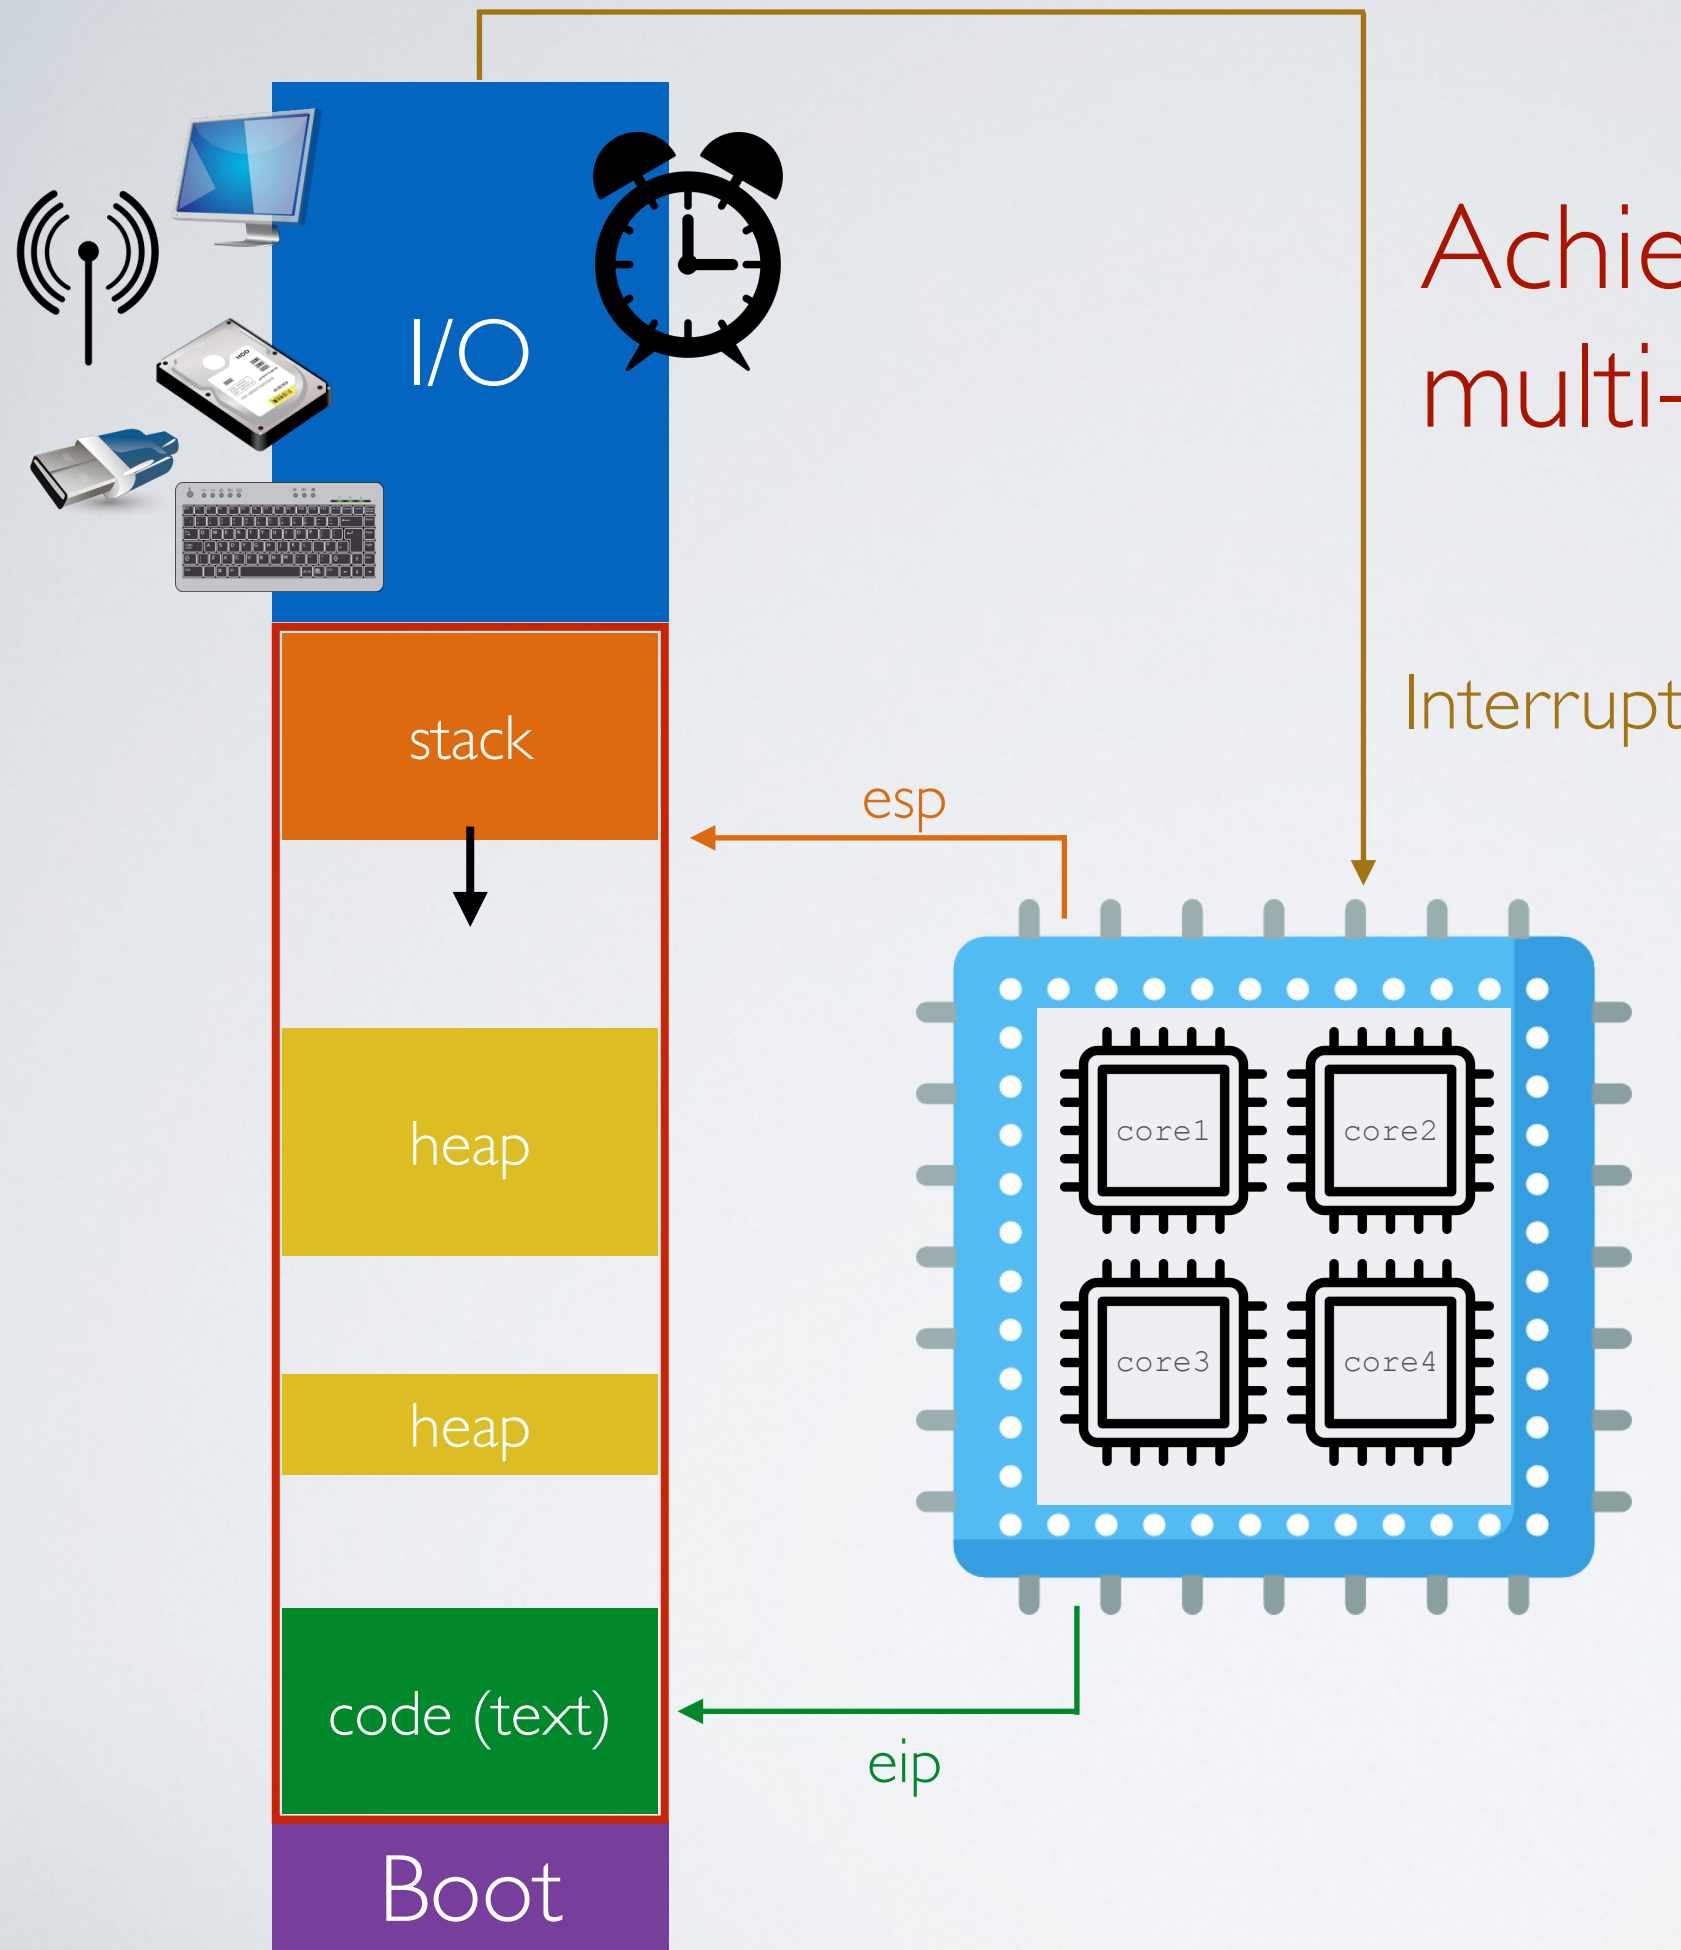
better CPU usage resulting in a faster execution overall  
(but not individually)

- ✓ **From the user perspective**

programs are executed in parallel

➡ But it requires scheduling, synchronization and some protection mechanisms





Achieving parallelism with multi-core processors

# Other problems that we are going to address during the semester

- **Scheduling**

Decide which process to execute when several are ready to be run

- **Synchronization**

Manage concurrent access to resources using semaphores, locks, monitors

- **Communication**

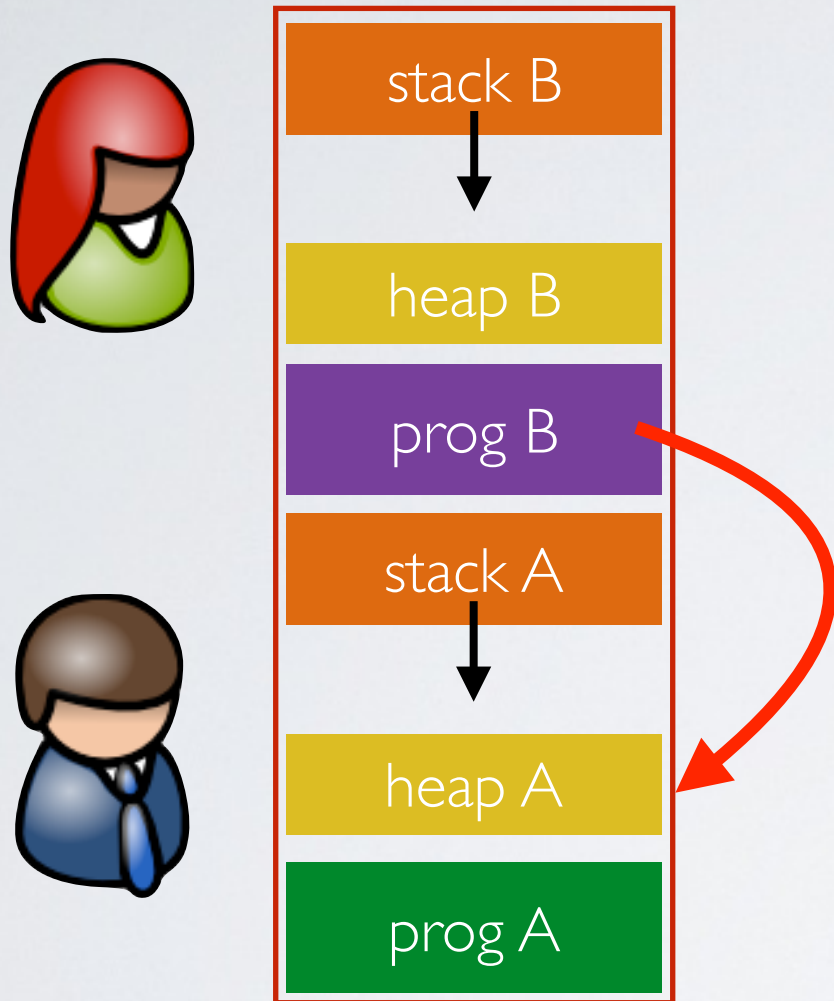
Exchange messages between processes using IPC (sockets & signals)

- **Threads**

Lightweight concurrency within a process

The need for **user spaces**

# An old problem from older constraints



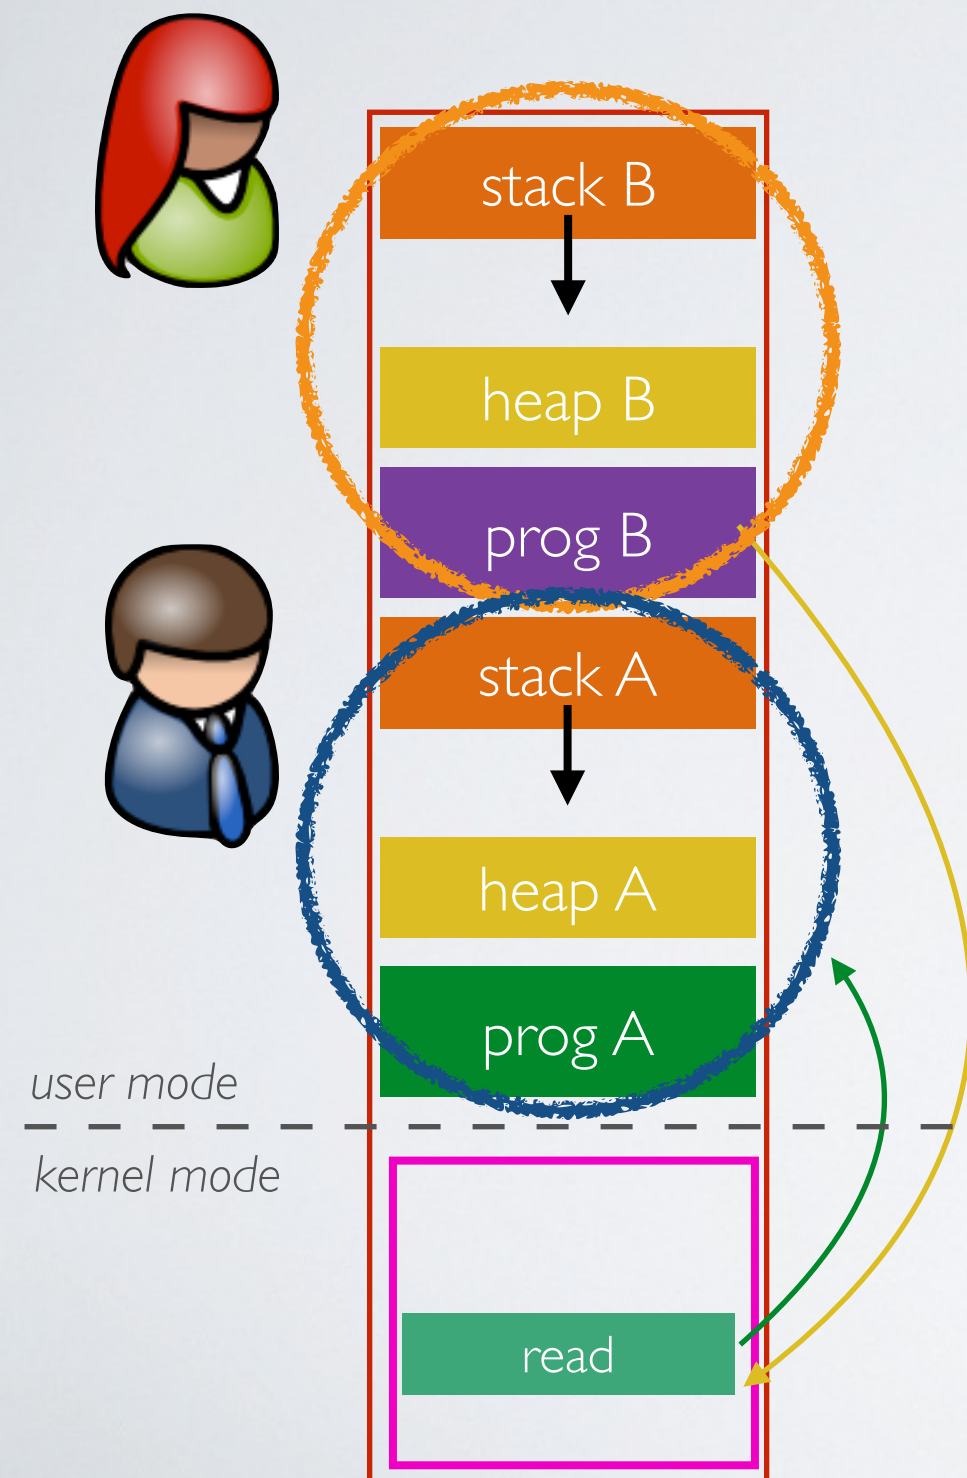
One computer and many users

**Problem:** what prevents Alice's from reading Bob's data?

- or start/stop any programs?
- or access any file on the filesystem
- or use any I/O device?
- or change the system configuration?
- or reboot the machine?



# Definition of the user space



**principle 1:** user have full privileges with their own user space

**principle 2:** every access to another user space must go through the kernel via system calls (complete mediation)

**principle 3:** system calls can be allowed or denied based on the system security policy (access control)

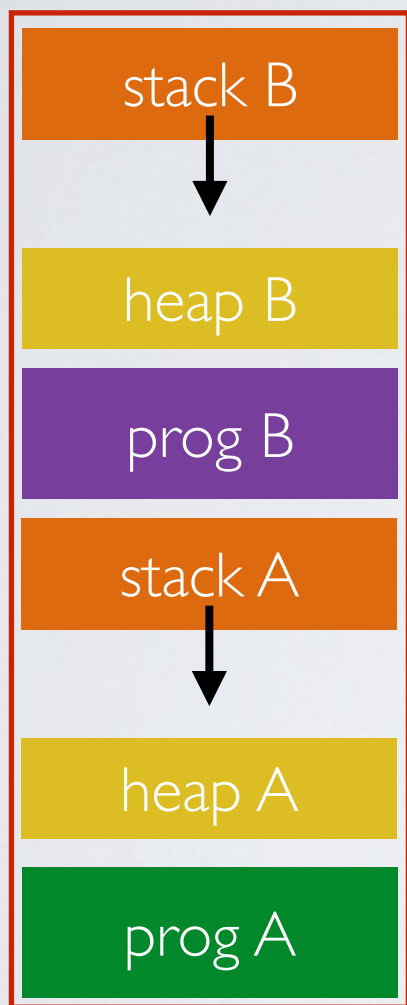


# Is multi-user paradigm obsolete?

- ➔ Most servers, personal computers, mobile and embedded systems have a single physical user
- ⦿ But not all programs are reliable nor trustworthy
- ✓ It is still a good model to provide **reliability** and **security**

The need for **virtual memory**

# The problem of managing the memory



How to make programs and execution contexts co-exists in memory?

- ✓ Placing multiple execution contexts (stack and heap) at random locations in memory is not a problem ...  
... well, as long as you have enough memory
- ⦿ However having programs placed at random locations is problematic

# Let's look at some C code and its binary

```
#include <stdio.h>

int foo(){
    printf("hello world!");
}

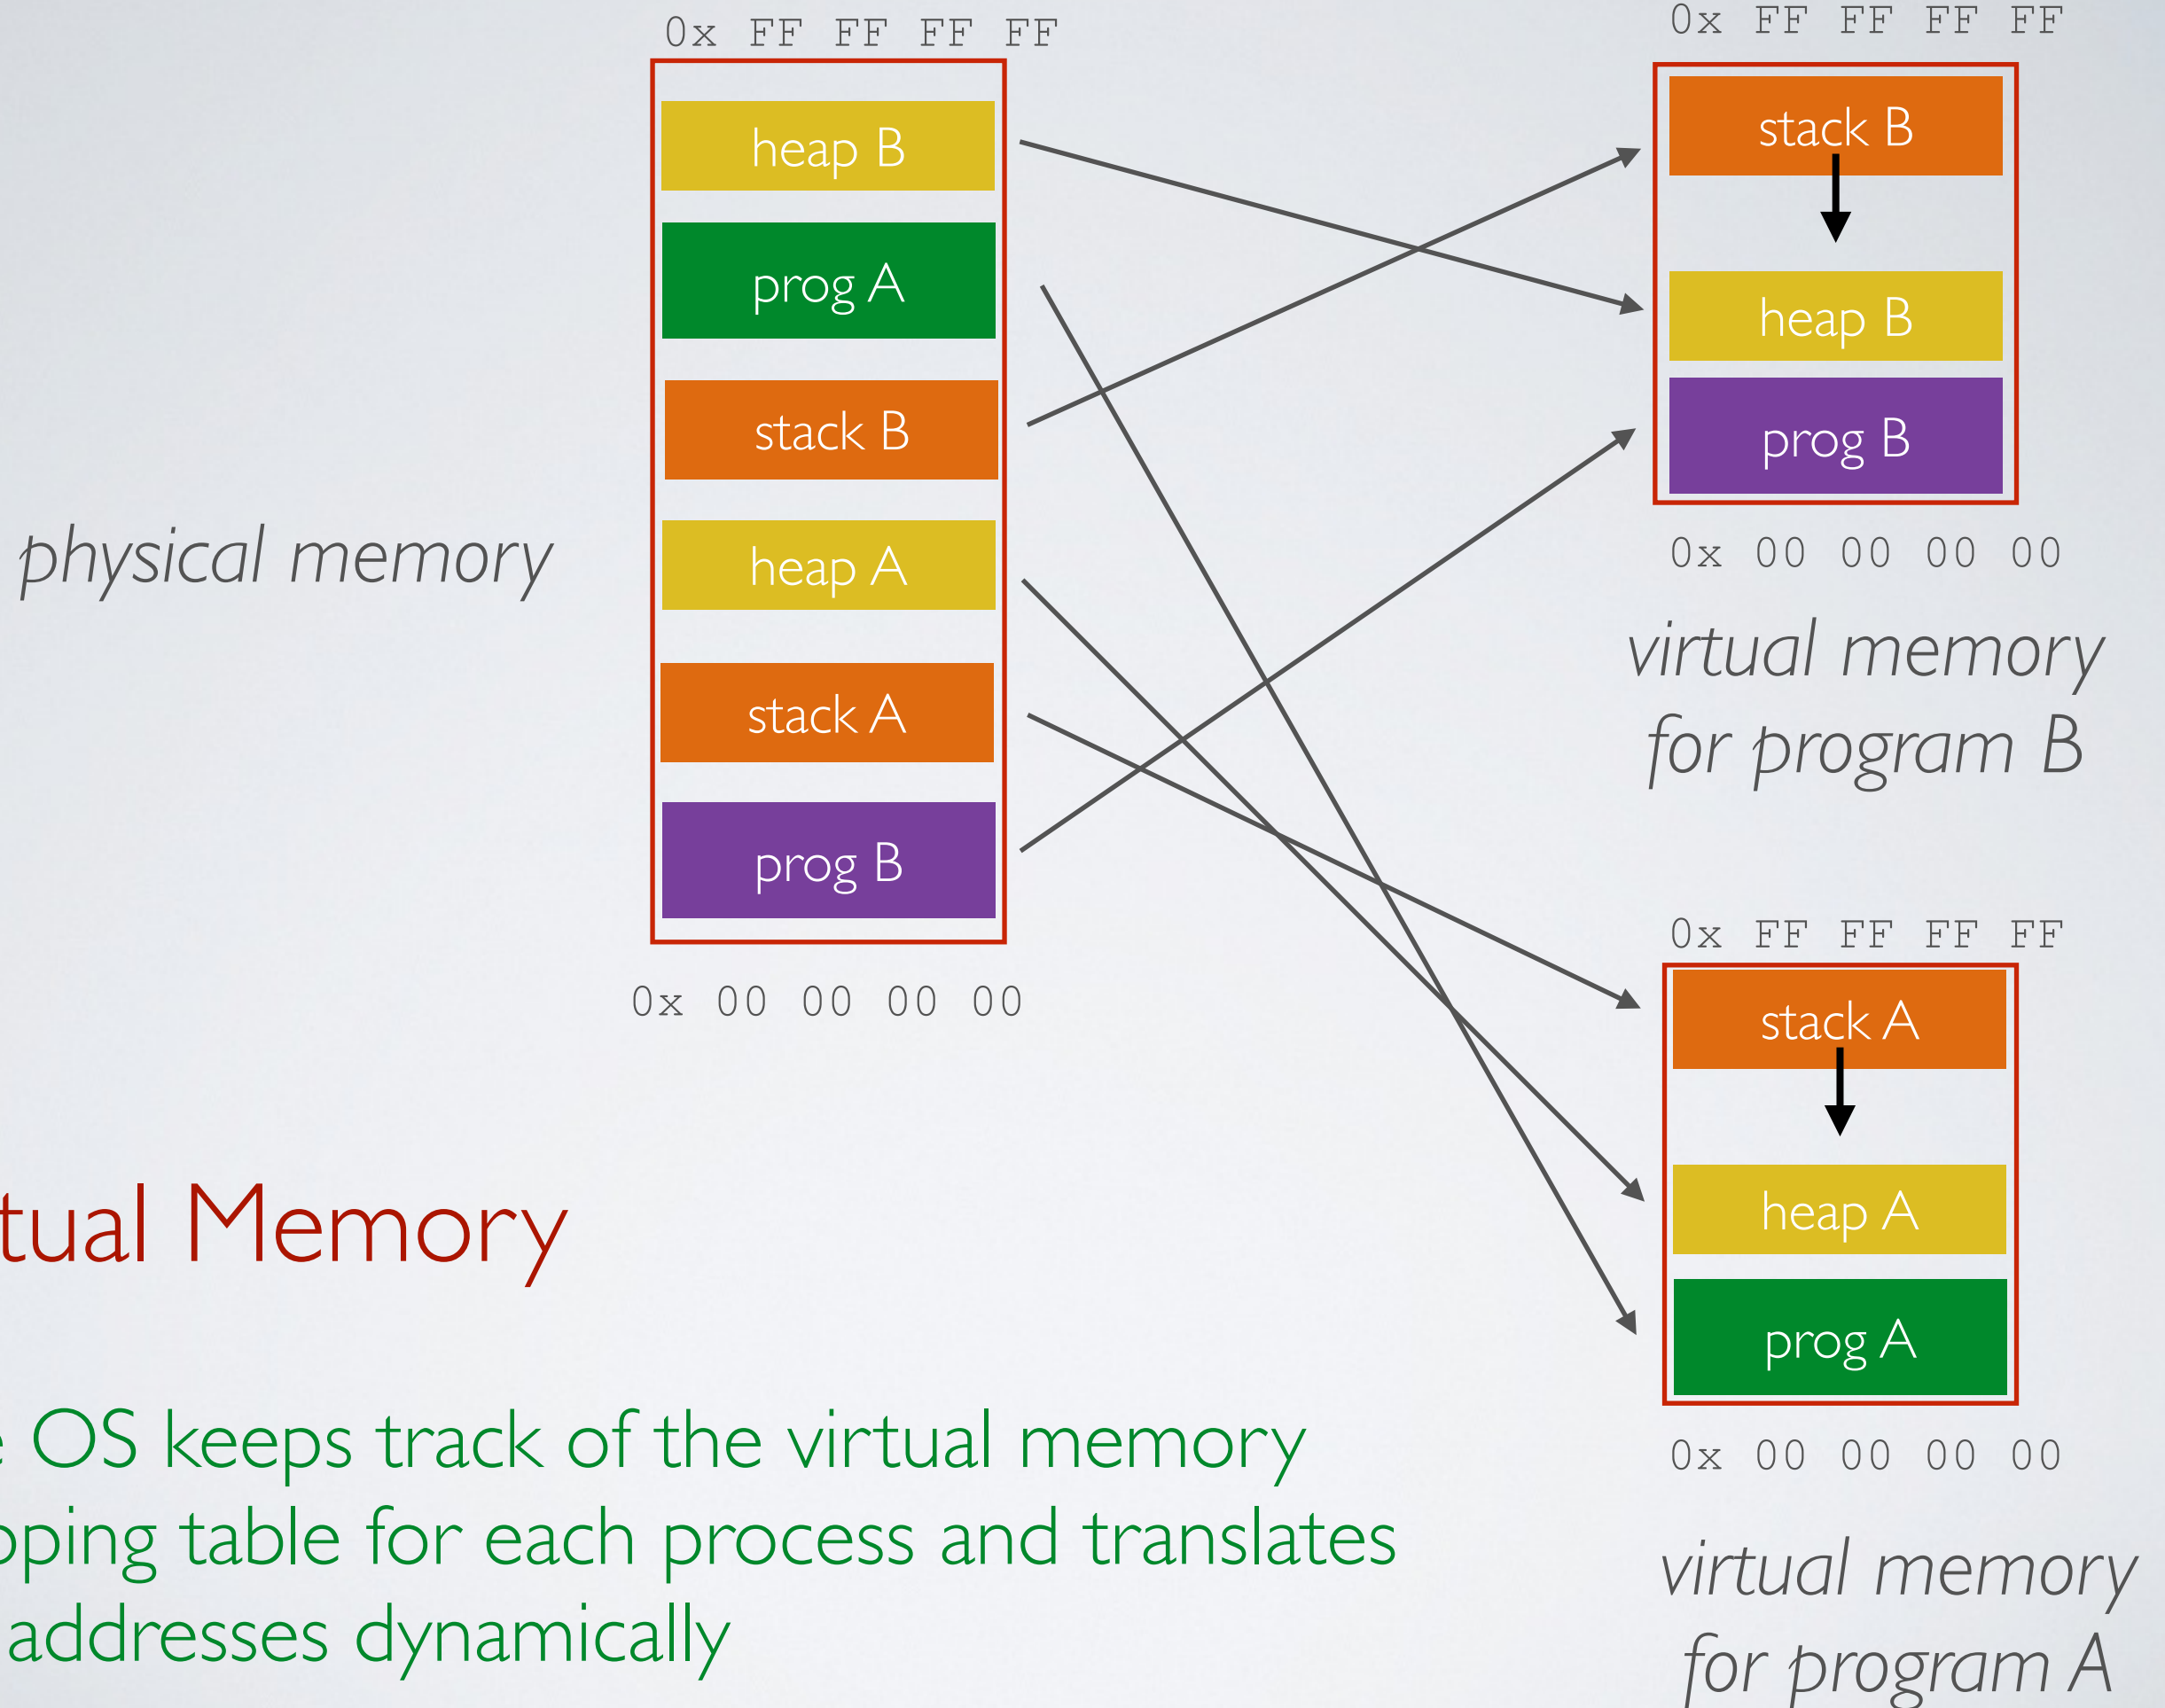
int main(int argc, char **argv){
    foo();
}
```

Since function addresses and others are hard-encoded in the binary, the program cannot be placed at random locations in memory

```
0804840b <foo>:
804840b: 55                push    ebp
804840c: 89 e5            mov     ebp,esp
804840e: 83 ec 08         sub     esp,0x8
8048411: 83 ec 0c         sub     esp,0xc
8048414: 68 d0 84 04 08   push    0x80484d0
8048419: e8 c2 fe ff ff   call    80482e0 <printf@plt>
804841e: 83 c4 10         add     esp,0x10
8048421: 90              nop
8048422: c9              leave
8048423: c3              ret

08048424 <main>:
8048424: 8d 4c 24 04      lea     ecx,[esp+0x4]
8048428: 83 e4 f0         and     esp,0xffffffff0
804842b: ff 71 fc         push    DWORD PTR [ecx-0x4]
804842e: 55              push    ebp
804842f: 89 e5            mov     ebp,esp
8048431: 51              push    ecx
8048432: 83 ec 04         sub     esp,0x4
8048435: e8 d1 ff ff ff   call    804840b <foo>
804843a: b8 00 00 00 00   mov     eax,0x0
804843f: 83 c4 04         add     esp,0x4
8048442: 59              pop     ecx
8048443: 5d              pop     ebp
8048444: 8d 61 fc         lea     esp,[ecx-0x4]
8048447: c3              ret
8048448: 66 90           xchg    ax,ax
804844a: 66 90           xchg    ax,ax
804844c: 66 90           xchg    ax,ax
804844e: 66 90           xchg    ax,ax
```







# Another problem


What if we run out of memory because of too many concurrent programs?

✓ Swap memory  
move some data to the disk

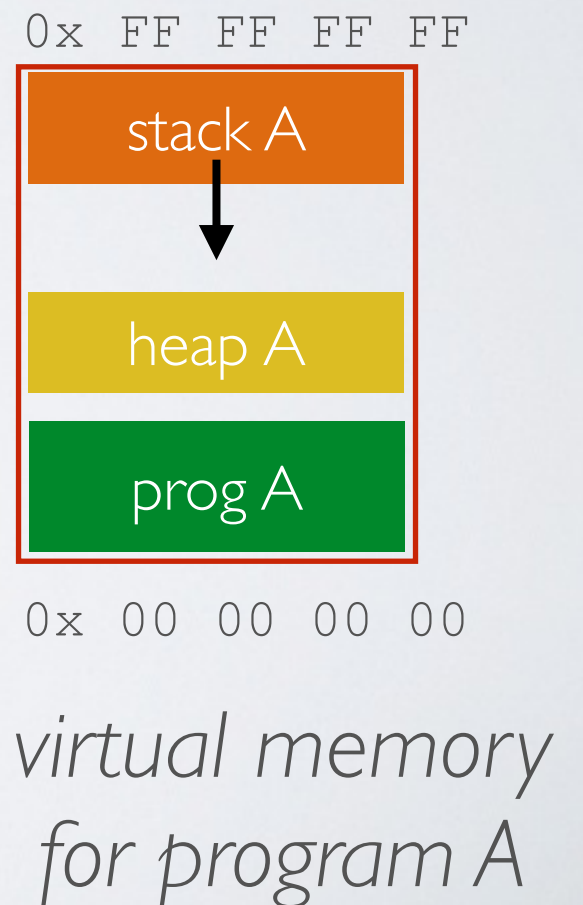

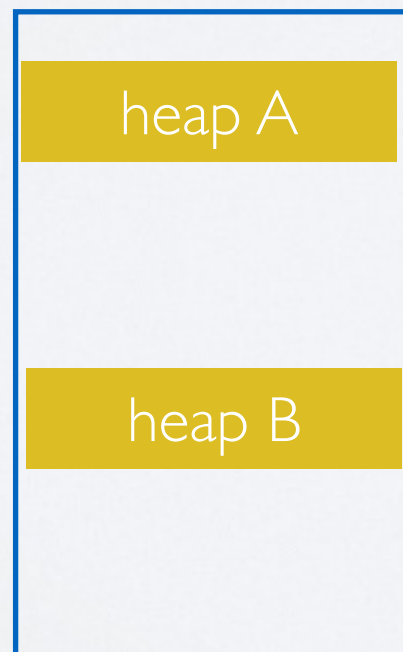
➡ Managing memory becomes very complex  
but necessary

# Swap

*physical memory*

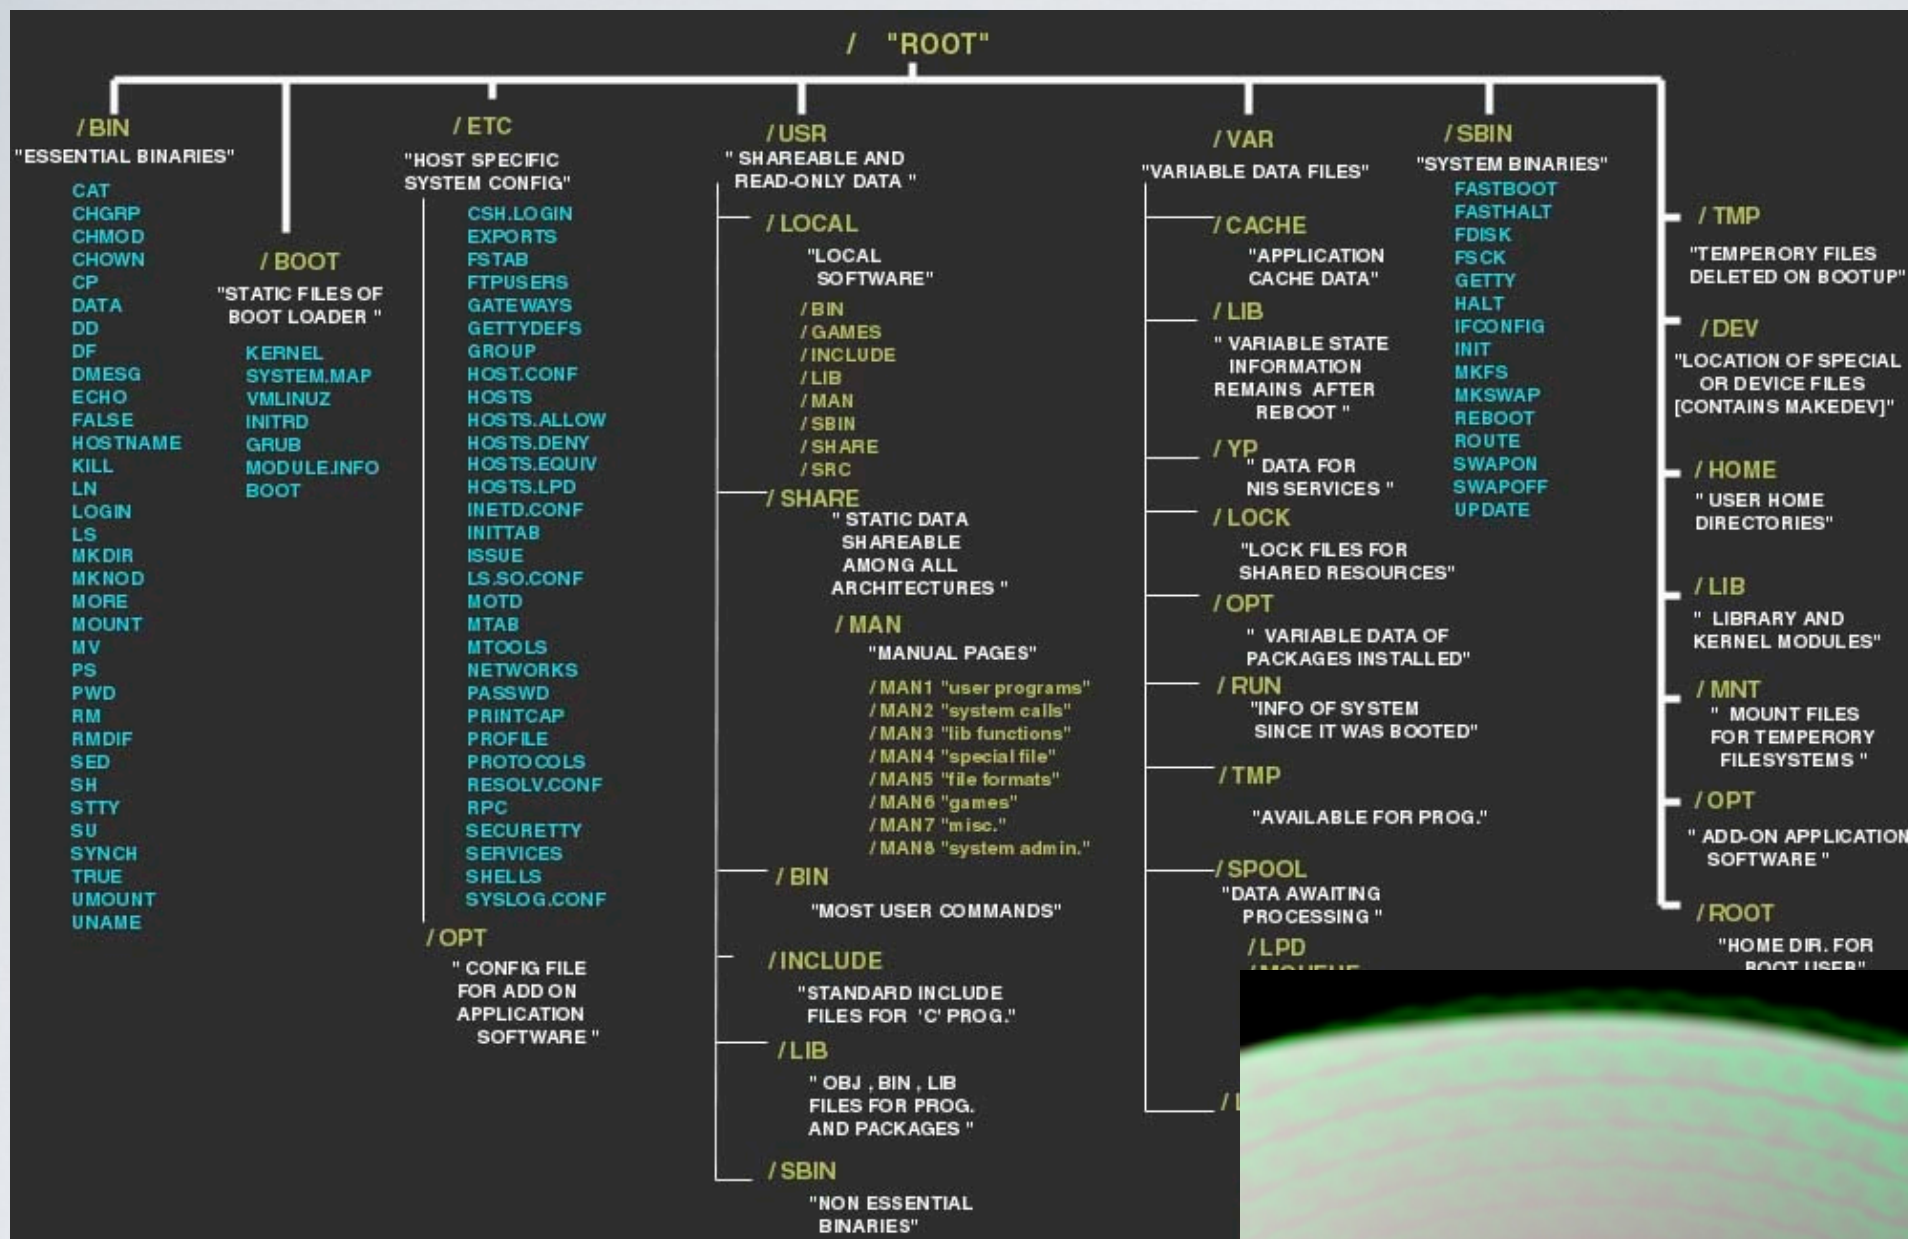


*hard drive*



The need for a **file system**





Files and Directories

versus

Reality





So, what is an **operating system**?

# Operating System

- ➔ In a nutshell, an OS manages hardware and runs programs
  - creates and manages processes
  - manages access to the memory (including RAM and I/O)
  - manages files and directories of the filesystem on disk(s)
  - enforces protection mechanisms for reliability and security
  - enables inter-process communication

# For next week

- Read the book
- Read Pintos documentations  
(0-Introduction and A-Reference Guide)
- Work on Project 0  
(to be finalized in the next couple of days)